

Applied Machine Learning and Predictive Modelling I: Modelling Stroke Data

Authors: Larissa Eisele, Fabian Lüthard, Yves Maillard

Module: Applied Machine Learning and Predictive Modelling I

Submitted on 10th of June, 2022

SUPERVISOR: MATTEO TANADINI AND DANIEL MEISTER

Table of Contents

1	Introduction	1
2	Importing Data	1
2.1	Summaries	1
2.2	Scatterplots	2
2.3	Boxplots	6
3	Methodology	12
4	Linear Model	12
5	Generalised Linear Model with family set to Poisson	13
6	Generalised Linear Model with family set to Binomial	14
7	Generalised Additive Model	14
8	Neural Network	14
9	Support Vector Machine	14
10	solve an optimisation problem	14

1 Introduction

Use case: We are a smart watch manufacturer working on a new feature for stroke prevention. We are going to analyze survey data that we plan to ask our users, complementing it with HR (Heart Rate) and CGM (Continuous Glucose Monitoring) data that our product already measures. We hope that our feature can prevent serious health issues and motivate our users to adopt healthier lifestyles.

2 Importing Data

```
stroke_data <- read_csv('./data/healthcare-dataset-stroke-data.csv')

## Rows: 5110 Columns: 12

## -- Column specification -----
## Delimiter: ","
## chr (6): gender, ever_married, work_type, Residence_type, bmi, smoking_status
## dbl (6): id, age, hypertension, heart_disease, avg_glucose_level, stroke

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
stroke_data$bmi <- as.numeric(stroke_data$bmi)

## Warning: NAs introduced by coercion
stroke_data$bmi[is.na(stroke_data$bmi)] <- mean(stroke_data$bmi, na.rm = TRUE)
stroke_data$stroke <- as.factor(stroke_data$stroke)

stroke_data

## # A tibble: 5,110 x 12
##       id gender   age hypertension heart_disease ever_married work_type
##   <dbl> <chr>   <dbl>         <dbl>         <dbl> <chr>         <chr>
## 1  9046 Male     67             0             1 Yes         Private
## 2  51676 Female  61             0             0 Yes         Self-employed
## 3  31112 Male    80             0             1 Yes         Private
## 4  60182 Female  49             0             0 Yes         Private
## 5   1665 Female  79             1             0 Yes         Self-employed
## 6  56669 Male    81             0             0 Yes         Private
## 7  53882 Male    74             1             1 Yes         Private
## 8  10434 Female  69             0             0 No          Private
## 9  27419 Female  59             0             0 Yes         Private
## 10 60491 Female  78             0             0 Yes         Private
## # ... with 5,100 more rows, and 5 more variables: Residence_type <chr>,
## #   avg_glucose_level <dbl>, bmi <dbl>, smoking_status <chr>, stroke <fct>
```

2.1 Summaries

```
#testing the effect of non smokers and smokers
count_by_smoke_status <- stroke_data %>%
  select(smoking_status, stroke) %>%
  group_by(smoking_status, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'smoking_status'. You can override using the `.groups` argument.

```
#testing the effect of work type
count_by_work_type <- stroke_data %>%
  select(work_type, stroke) %>%
  group_by(work_type, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'work_type'. You can override using the `.groups` argument.

```
# testing the effects of residence type
count_by_Residence_type <- stroke_data %>%
  select(Residence_type, stroke) %>%
  group_by(Residence_type, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'Residence_type'. You can override using the `.groups` argument.

```
# testing the effects of gender
count_by_gender <- stroke_data %>%
  select(gender, stroke) %>%
  group_by(gender, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'gender'. You can override using the `.groups` argument.

```
# testing the effects of hypertension
count_by_hypertension <- stroke_data %>%
  select(hypertension, stroke) %>%
  group_by(hypertension, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'hypertension'. You can override using the `.groups` argument.

```
# testing the effects of heart disease
count_by_heart_disease <- stroke_data %>%
  select(heart_disease, stroke) %>%
  group_by(heart_disease, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'heart_disease'. You can override using the `.groups` argument.


```
# testing the effects of marriage status
count_by_marriage <- stroke_data %>%
  select(ever_married, stroke) %>%
  group_by(ever_married, stroke) %>%
  summarise( N = n())
```

`summarise()` has grouped output by 'ever_married'. You can override using the `.groups` argument.

2.2 Scatterplots

```
stroke_data %>%
  ggplot(mapping = aes(x = bmi, y = age, color = stroke)) +
  geom_point() +
  geom_smooth(method = 'lm')
```


`geom_smooth()` using formula 'y ~ x'



```
stroke_data %>%  
  filter(stroke == 1) %>%  
  ggplot(mapping = aes(x = bmi, y = age, color = smoking_status)) +  
  geom_point(method = 'lm') +  
  geom_smooth(method = 'lm')
```


```
## Warning: Ignoring unknown parameters: method
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
stroke_data %>%
  ggplot(mapping = aes(x = bmi, y = avg_glucose_level, color = stroke)) +
  geom_point() +
  geom_smooth(method = 'lm')
```


```
## `geom_smooth()` using formula 'y ~ x'
```



```
stroke_data %>%  
  filter(stroke == 1) %>%  
  ggplot(mapping = aes(x = bmi, y = avg_glucose_level, color = stroke)) +  
  geom_point(method = 'lm') +  
  geom_smooth(method = 'lm')
```


```
## Warning: Ignoring unknown parameters: method
```

```
## `geom_smooth()` using formula 'y ~ x'
```




2.3 Boxplots

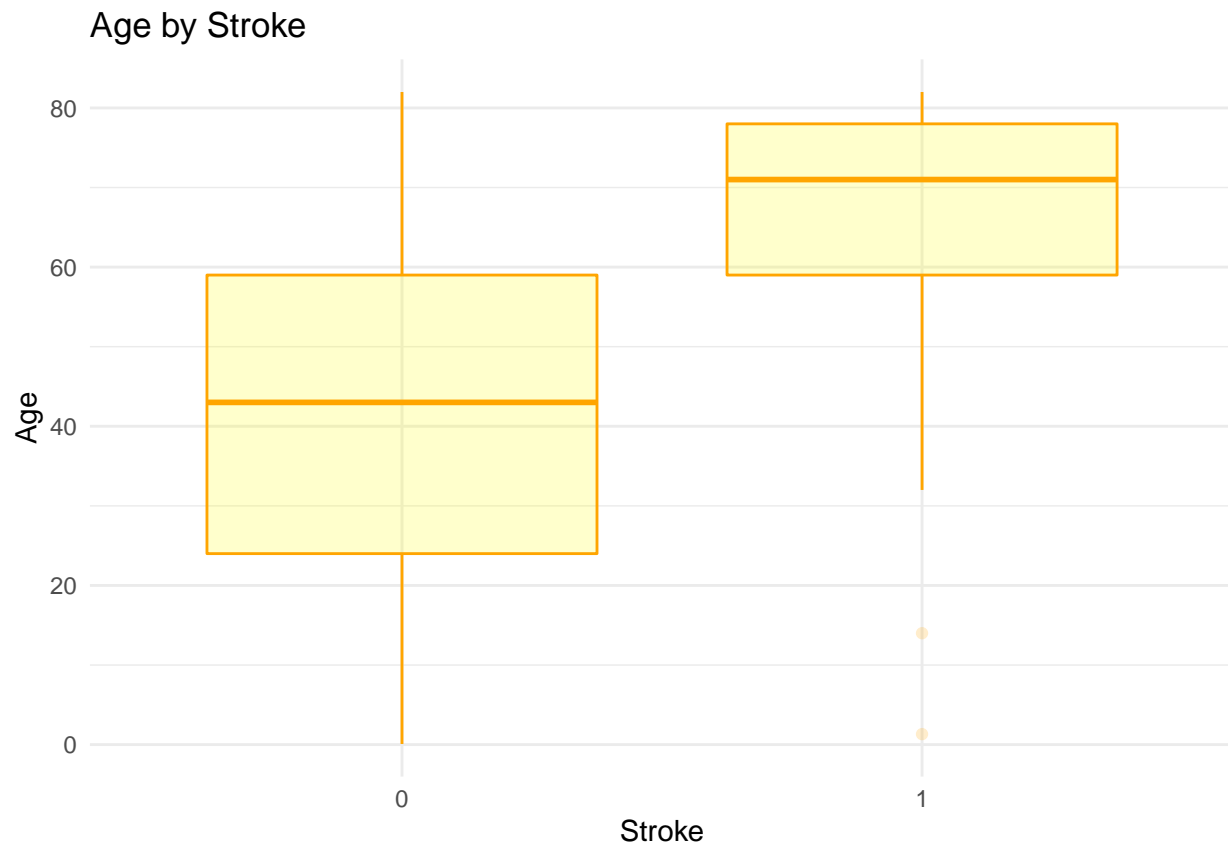
```
#Boxplot on avg. glucose_level  
ggplot(stroke_data, aes(y = avg_glucose_level)) +  
  geom_boxplot()
```


```
#Boxplot on Age  
ggplot(stroke_data, aes(y = age)) +  
  geom_boxplot()
```




```
stroke_by_age <- stroke_data %>%  
# dplyr::filter(stroke == 1) %>%  
ggplot(aes(x = stroke,  
           y = age)) +  
  geom_boxplot(color="orange", fill="yellow", alpha=0.2) +  
  ggtitle("Age by Stroke") +  
  xlab("Stroke") + ylab("Age") +  
  theme_minimal() + theme(axis.text.x = element_text(angle = 0))  
  
stroke_by_age
```



```
#Boxplot on BMI  
ggplot(stroke_data, aes(y = bmi)) +  
geom_boxplot()
```

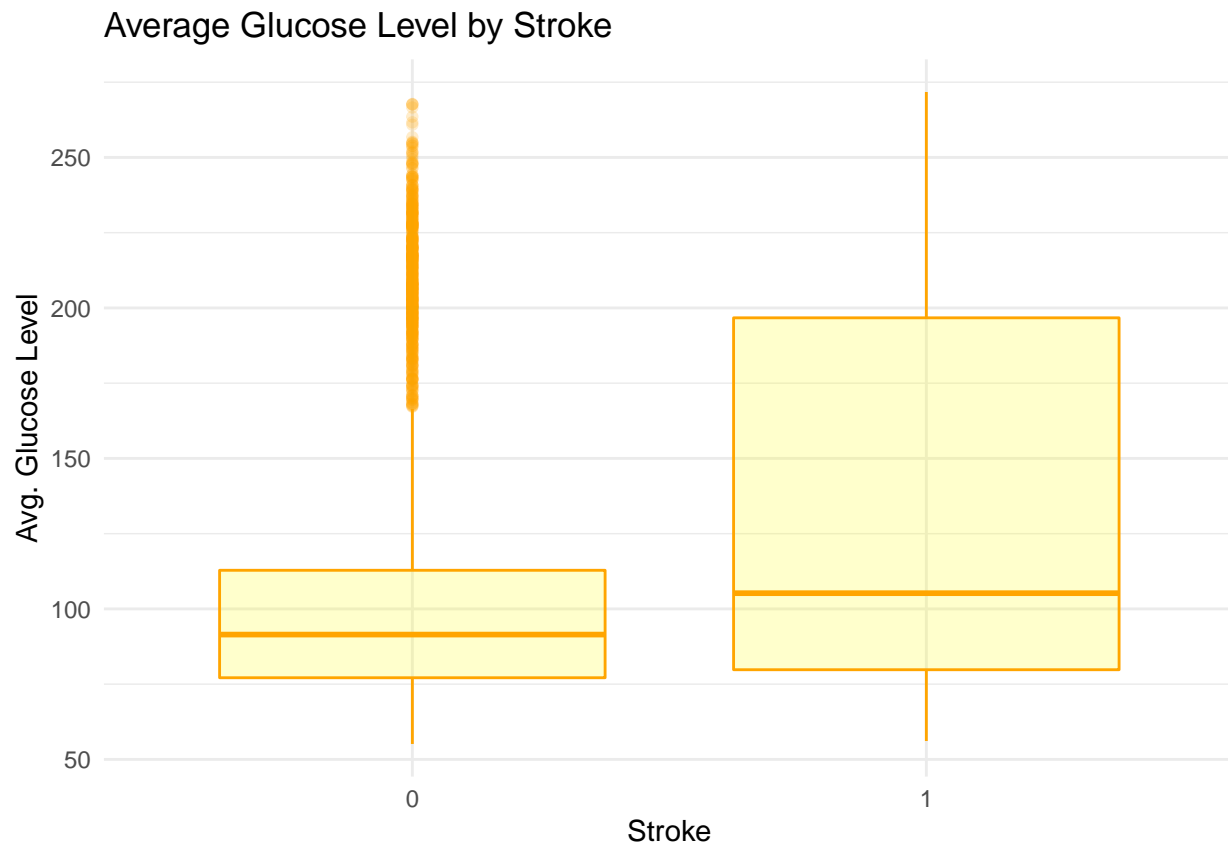


```
stroke_by_bmi <- stroke_data %>%  
# dplyr::filter(stroke == 1) %>%  
ggplot(aes(x = stroke,  
           y = bmi)) +  
  geom_boxplot(color="orange", fill="yellow", alpha=0.2) +  
  ggtitle("BMI by Stroke") +  
  xlab("Stroke") + ylab("BMI") +  
  theme_minimal() + theme(axis.text.x = element_text(angle = 0))  
stroke_by_bmi
```



```
#boxplot by avg_glucose_level and stroke
stroke_by_avg_glucose_level <- stroke_data %>%
# dplyr::filter(stroke == 1) %>%
ggplot(aes(x = stroke,
            y = avg_glucose_level)) +
  geom_boxplot(color="orange", fill="yellow", alpha=0.2) +
  ggtitle("Average Glucose Level by Stroke") +
  xlab("Stroke") + ylab("Avg. Glucose Level") +
  theme_minimal() + theme(axis.text.x = element_text(angle = 0))


stroke_by_avg_glucose_level
```



3 Methodology

4 Linear Model

```
# qplot -> age/glucose level per stroke
library(ggplot2)
qplot(y = age, x = avg_glucose_level,
      data = stroke_data,
      facets = ~ stroke)
```



```
# fitting models for simple regression model
lm.stroke <- lm(age ~ avg_glucose_level, data = stroke_data)
summary(lm.stroke)

##
## Call:
## lm(formula = age ~ avg_glucose_level, data = stroke_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -53.368 -16.775   1.115  16.626  44.648
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    30.60224    0.783098   39.08  <2e-16 ***
## avg_glucose_level  0.118932    0.006786   17.53  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.96 on 5108 degrees of freedom
## Multiple R-squared:  0.05673,    Adjusted R-squared:  0.05654
## F-statistic: 307.2 on 1 and 5108 DF,  p-value: < 2.2e-16
##
```

5 Generalised Linear Model with family set to Poisson

- 6 Generalised Linear Model with family set to Binomial**
- 7 Generalised Additive Model**
- 8 Neural Network**
- 9 Support Vector Machine**
- 10 solve an optimisation problem**