

Applied Machine Learning and Predictive Modelling I: Modelling Stroke Data

Authors: Larissa Eisele, Fabian Lüthard, Yves Maillard

Module: Applied Machine Learning and Predictive Modelling I

Submitted on 10th of June, 2022

SUPERVISOR: MATTEO TANADINI AND DANIEL MEISTER

Table of Contents

1	Introduction	1
2	Importing Data	1
2.1	Summaries	2
2.2	Scatterplots	3
2.3	Boxplots	7
3	Methodology	13
4	Linear Models	14
5	Generalised Linear Model with family set to Poisson	25
6	Generalised Linear Model with family set to Binomial	28
7	Generalised Additive Model	41
8	Neural Network Yves	42
9	Support Vector Machine (Larissa)	42
10	OPTIONAL solve an optimisation problem	43

1 Introduction

Use case: We are a smart watch manufacturer working on a new feature for stroke prevention. We are going to analyze survey data that we plan to ask our users, complementing it with HR (Heart Rate) and CGM (Continuous Glucose Monitoring) data that our product already measures. We hope that our feature can prevent serious health issues and motivate our users to adopt healthier lifestyles.

2 Importing Data

```
stroke_data <- read_csv('./data/healthcare-dataset-stroke-data.csv')

## Rows: 5110 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (6): gender, ever_married, work_type, Residence_type, bmi, smoking_status
## dbl (6): id, age, hypertension, heart_disease, avg_glucose_level, stroke
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

s <- spec(stroke_data)
s

## cols(
##   id = col_double(),
##   gender = col_character(),
##   age = col_double(),
##   hypertension = col_double(),
##   heart_disease = col_double(),
##   ever_married = col_character(),
##   work_type = col_character(),
##   Residence_type = col_character(),
##   avg_glucose_level = col_double(),
##   bmi = col_character(),
##   smoking_status = col_character(),
##   stroke = col_double()
## )

cols_condense(s)

## cols(
##   .default = col_character(),
##   id = col_double(),
##   age = col_double(),
##   hypertension = col_double(),
##   heart_disease = col_double(),
##   avg_glucose_level = col_double(),
##   stroke = col_double()
## )

stroke_data$bmi <- as.numeric(stroke_data$bmi)

## Warning: NAs introduced by coercion

stroke_data$bmi[is.na(stroke_data$bmi)] <- mean(stroke_data$bmi, na.rm = TRUE)
stroke_data$age <- as.integer(stroke_data$age)
```

```
stroke_data$smoking_status <- as.factor(stroke_data$smoking_status)
stroke_data$work_type <- as.factor(stroke_data$work_type)
stroke_data$gender <- as.factor(stroke_data$gender)
stroke_data$ever_married <- as.factor(stroke_data$ever_married)
stroke_data$Residence_type <- as.factor(stroke_data$Residence_type)
stroke_data$stroke_num <- as.numeric(stroke_data$stroke)
stroke_data$stroke <- as.factor(stroke_data$stroke)
```

```
stroke_data
```

```
## # A tibble: 5,110 x 13
##       id gender   age hypertension heart_disease ever_married work_type
##   <dbl> <fct> <int>         <dbl>         <dbl> <fct>         <fct>
## 1  9046 Male    67           0           1 Yes      Private
## 2 51676 Female  61           0           0 Yes      Self-employed
## 3 31112 Male    80           0           1 Yes      Private
## 4 60182 Female  49           0           0 Yes      Private
## 5  1665 Female  79           1           0 Yes      Self-employed
## 6 56669 Male    81           0           0 Yes      Private
## 7 53882 Male    74           1           1 Yes      Private
## 8 10434 Female  69           0           0 No       Private
## 9 27419 Female  59           0           0 Yes      Private
##10 60491 Female  78           0           0 Yes      Private
## # ... with 5,100 more rows, and 6 more variables: Residence_type <fct>,
## #   avg_glucose_level <dbl>, bmi <dbl>, smoking_status <fct>, stroke <fct>,
## #   stroke_num <dbl>
```

2.1 Summaries

```
#testing the effect of non smokers and smokers
```

```
count_by_smoke_status <- stroke_data %>%
  select(smoking_status, stroke) %>%
  group_by(smoking_status, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'smoking_status'. You can override using the
## `.groups` argument.
```

```
#testing the effect of work type
```

```
count_by_work_type <- stroke_data %>%
  select(work_type, stroke) %>%
  group_by(work_type, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'work_type'. You can override using the
## `.groups` argument.
```

```
# testing the effects of residence type
```

```
count_by_Residence_type <- stroke_data %>%
  select(Residence_type, stroke) %>%
  group_by(Residence_type, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'Residence_type'. You can override using the
## `.groups` argument.
```

```
# testing the effects of gender
count_by_gender <- stroke_data %>%
  select(gender, stroke) %>%
  group_by(gender, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'gender'. You can override using the
## `.groups` argument.
```

```
# testing the effects of hypertension
count_by_hypertension <- stroke_data %>%
  select(hypertension, stroke) %>%
  group_by(hypertension, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'hypertension'. You can override using the
## `.groups` argument.
```

```
# testing the effects of heart disease
count_by_heart_disease <- stroke_data %>%
  select(heart_disease, stroke) %>%
  group_by(heart_disease, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'heart_disease'. You can override using the
## `.groups` argument.
```


```
# testing the effects of marriage status
count_by_marriage <- stroke_data %>%
  select(ever_married, stroke) %>%
  group_by(ever_married, stroke) %>%
  summarise( N = n())
```

```
## `summarise()` has grouped output by 'ever_married'. You can override using the
## `.groups` argument.
```

2.2 Scatterplots

```
stroke_data %>%
  ggplot(mapping = aes(x = bmi, y = age, color = stroke)) +
  geom_point() +
  geom_smooth(method = 'lm')
```


```
## `geom_smooth()` using formula 'y ~ x'
```



```
stroke_data %>%  
  filter(stroke == 1) %>%  
  ggplot(mapping = aes(x = bmi, y = age, color = smoking_status)) +  
  geom_point(method = 'lm') +  
  geom_smooth(method = 'lm')
```


```
## Warning: Ignoring unknown parameters: method
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
stroke_data %>%
  ggplot(mapping = aes(x = bmi, y = avg_glucose_level, color = stroke)) +
  geom_point() +
  geom_smooth(method = 'lm')
```


```
## `geom_smooth()` using formula 'y ~ x'
```



```
stroke_data %>%  
  filter(stroke == 1) %>%  
  ggplot(mapping = aes(x = bmi, y = avg_glucose_level, color = stroke)) +  
  geom_point(method = 'lm') +  
  geom_smooth(method = 'lm')
```


```
## Warning: Ignoring unknown parameters: method
```

```
## `geom_smooth()` using formula 'y ~ x'
```





2.3 Boxplots


```
#Boxplot on avg. glucose_level  
ggplot(stroke_data, aes(y = avg_glucose_level)) +  
  geom_boxplot()
```




```
#Boxplot on Age  
ggplot(stroke_data, aes(y = age)) +  
  geom_boxplot()
```




```
stroke_by_age <- stroke_data %>%  
# dplyr::filter(stroke == 1) %>%  
ggplot(aes(x = stroke,  
           y = age)) +  
  geom_boxplot() +  
  labs(title = "Boxplot on no-Stroke (0) and Stroke (1) by Age",  
       x = "No-stroke (0) and Stroke (1)",  
       y = "Age",  
       color = "green")  
  
stroke_by_age
```



```
#Boxplot on BMI  
ggplot(stroke_data, aes(y = bmi)) +  
  geom_boxplot()
```




```
stroke_by_bmi <- stroke_data %>%
# dplyr::filter(stroke == 1) %>%
ggplot(aes(x = stroke,
            y = bmi)) +
  geom_boxplot(color="orange", fill="yellow", alpha=0.2) +
  ggtitle("BMI by Stroke") +
  xlab("Stroke") + ylab("BMI") +
  theme_minimal() + theme(axis.text.x = element_text(angle = 0))
stroke_by_bmi
```



```
#boxplot by avg_glucose_level and stroke
stroke_by_avg_glucose_level <- stroke_data %>%
# dplyr::filter(stroke == 1) %>%
ggplot(aes(x = stroke,
            y = avg_glucose_level)) +
  geom_boxplot(color="orange", fill="yellow", alpha=0.2) +
  ggtitle("Average Glucose Level by Stroke") +
  xlab("Stroke") + ylab("Avg. Glucose Level") +
  theme_minimal() + theme(axis.text.x = element_text(angle = 0))

stroke_by_avg_glucose_level
```



3 Methodology

```
set.seed(7406)
n=dim(stroke_data[1]) # number of observations in dataset
n_train=0.70*n # training set is 70%
flag = sort(sample(1:n, size=n_train, replace=FALSE))

## Warning in 1:n: numerical expression has 2 elements: only the first used
# Use df (all data points without ID column) df_train, and df_test
# Gender, hypertension, heart disease, ever married, work type, residence type, smoking status, and str
# This should allow for the best modeling options possible for our methods.
df_train = stroke_data[flag,]
df_test = stroke_data[-flag,]

head(df_test)

## # A tibble: 6 x 13
##   id gender  age hypertension heart_disease ever_married work_type
##   <dbl> <fct> <int>         <dbl>         <dbl> <fct>         <fct>
## 1 51676 Female   61             0             0 Yes      Self-employed
## 2 10434 Female   69             0             0 No       Private
## 3 60491 Female   78             0             0 Yes      Private
## 4 12095 Female   61             0             1 Yes      Govt_job
## 5 58202 Female   50             1             0 Yes      Self-employed
## 6 56112 Male    64             0             1 Yes      Private
## # ... with 6 more variables: Residence_type <fct>, avg_glucose_level <dbl>,
```

```
## #   bmi <dbl>, smoking_status <fct>, stroke <fct>, stroke_num <dbl>
head(df_train)

## # A tibble: 6 x 13
##       id gender   age hypertension heart_disease ever_married work_type
##   <dbl> <fct> <int>         <dbl>         <dbl> <fct>         <fct>
## 1  9046 Male    67             0             1 Yes         Private
## 2 31112 Male    80             0             1 Yes         Private
## 3 60182 Female  49             0             0 Yes         Private
## 4  1665 Female  79             1             0 Yes         Self-employed
## 5 56669 Male    81             0             0 Yes         Private
## 6 53882 Male    74             1             1 Yes         Private
## # ... with 6 more variables: Residence_type <fct>, avg_glucose_level <dbl>,
## #   bmi <dbl>, smoking_status <fct>, stroke <fct>, stroke_num <dbl>
```

4 Linear Models

```
set.seed(7406)
n=dim(stroke_data[1]) # number of observations in dataset
n_train=0.70*n # training set is 70%
flag = sort(sample(1:n, size=n_train, replace=FALSE))

## Warning in 1:n: numerical expression has 2 elements: only the first used
# Use df (all data points without ID column) df_train_linear, and df_test_linear
# Gender, hypertension, heart disease, ever married, work type, residence type, smoking status, and stroke
# This should allow for the best modeling options possible for our methods.
df_train_linear = stroke_data[flag,]
df_test_linear = stroke_data[-flag,]

head(df_test_linear)

## # A tibble: 6 x 13
##       id gender   age hypertension heart_disease ever_married work_type
##   <dbl> <fct> <int>         <dbl>         <dbl> <fct>         <fct>
## 1 51676 Female  61             0             0 Yes         Self-employed
## 2 10434 Female  69             0             0 No          Private
## 3 60491 Female  78             0             0 Yes         Private
## 4 12095 Female  61             0             1 Yes         Govt_job
## 5 58202 Female  50             1             0 Yes         Self-employed
## 6 56112 Male    64             0             1 Yes         Private
## # ... with 6 more variables: Residence_type <fct>, avg_glucose_level <dbl>,
## #   bmi <dbl>, smoking_status <fct>, stroke <fct>, stroke_num <dbl>

head(df_train_linear)


## # A tibble: 6 x 13
##       id gender   age hypertension heart_disease ever_married work_type
##   <dbl> <fct> <int>         <dbl>         <dbl> <fct>         <fct>
## 1  9046 Male    67             0             1 Yes         Private
## 2 31112 Male    80             0             1 Yes         Private
## 3 60182 Female  49             0             0 Yes         Private
## 4  1665 Female  79             1             0 Yes         Self-employed
## 5 56669 Male    81             0             0 Yes         Private
## 6 53882 Male    74             1             1 Yes         Private
```



```
## # ... with 6 more variables: Residence_type <fct>, avg_glucose_level <dbl>,
## #   bmi <dbl>, smoking_status <fct>, stroke <fct>, stroke_num <dbl>
```

Our aim is to predict strokes, and the stroke variable is always our response variable of interest. However, simple or multiple linear regression is not the tool of choice, as stroke is a binary / categorical response variable. Thus in the following we will nonetheless fit a linear regression model to our data with stroke as the response variable (and another example to show that we've understood the concept). We know of the following limitations for fitting a linear regression model to a binary response variable: - - -

```
# qplot -> age/glucose level per stroke
library(ggplot2)
qplot(y = age, x = avg_glucose_level,
      data = df_train_linear,
      facets = ~ stroke)
```





```
# fitting models for simple regression model
lm.stroke <- lm(stroke_num ~ gender + age + hypertension + heart_disease + ever_married + work_type + R
summary(lm.stroke)
```


```
##
## Call:
## lm(formula = stroke_num ~ gender + age + hypertension + heart_disease +
##     ever_married + work_type + Residence_type + avg_glucose_level +
##     bmi + smoking_status, data = df_train_linear)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.33328 -0.08333 -0.02421  0.00834  1.03042
##
```

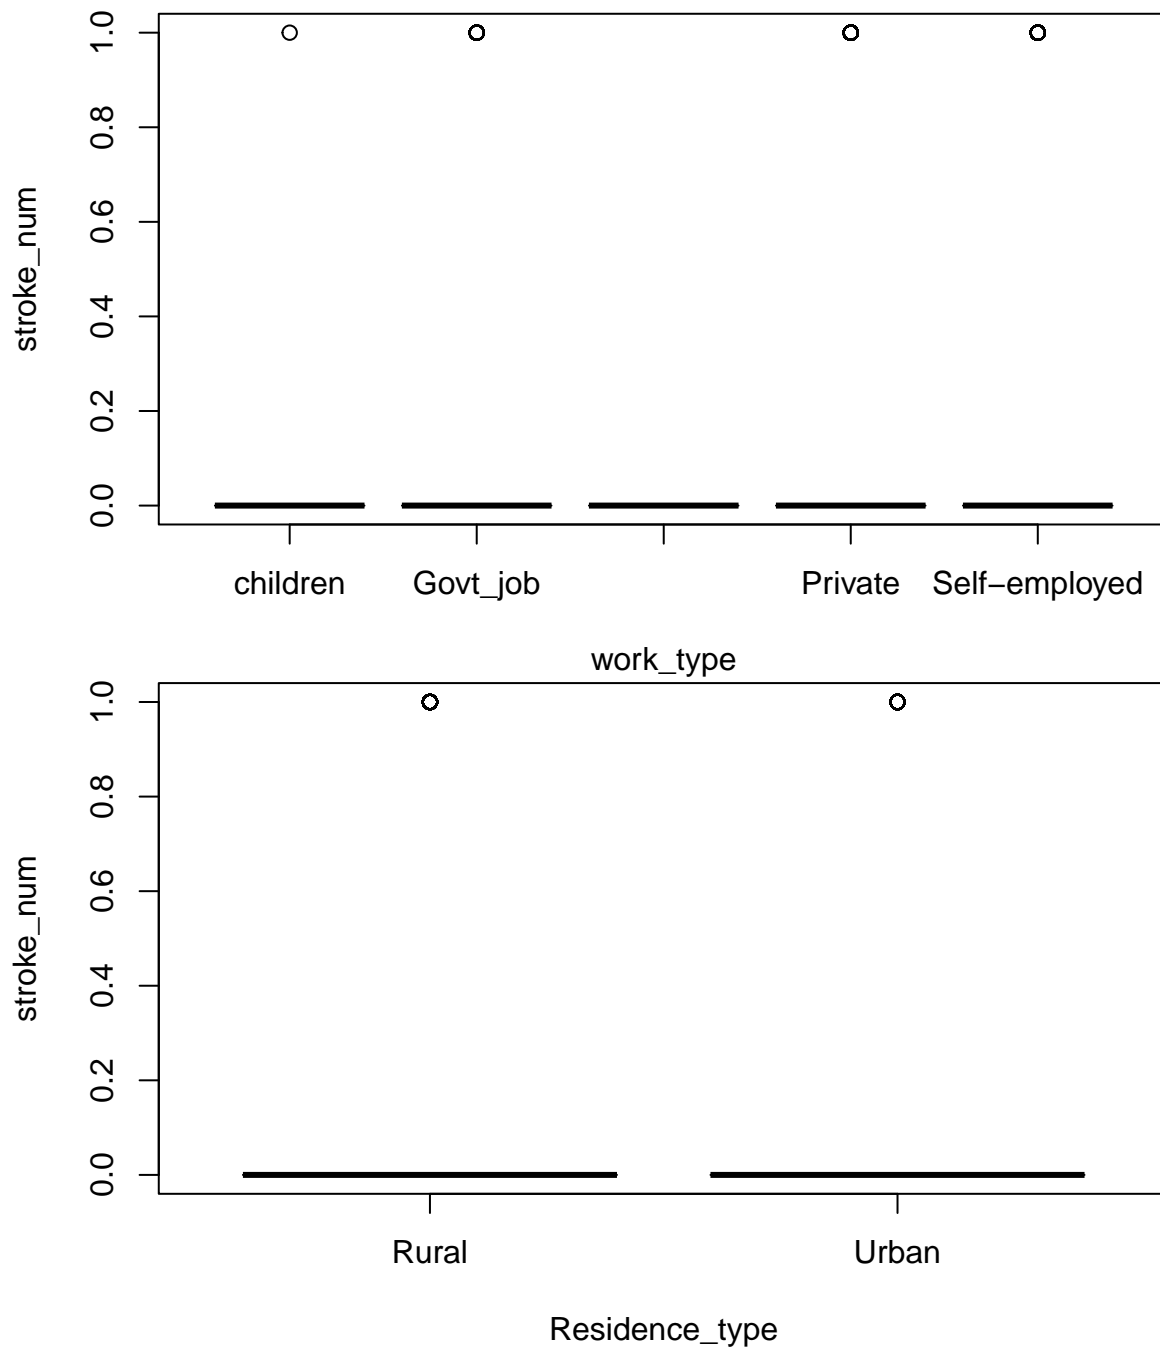
```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -3.812e-02  2.009e-02  -1.898 0.057814 .
## genderMale      4.560e-03  7.220e-03   0.632 0.527713
## genderOther    -2.698e-02  2.099e-01  -0.129 0.897733
## age            3.302e-03  2.617e-04  12.620 < 2e-16 ***
## hypertension   3.704e-02  1.210e-02   3.062 0.002218 **
## heart_disease   7.191e-02  1.598e-02   4.500 7.01e-06 ***
## ever_marriedYes -3.888e-02  1.045e-02  -3.721 0.000202 ***
## work_typeGovt_job -6.863e-02  1.854e-02  -3.701 0.000218 ***
## work_typeNever_worked -2.941e-02  5.372e-02  -0.548 0.584045
## work_typePrivate -5.075e-02  1.558e-02  -3.256 0.001139 **
## work_typeSelf-employed -8.413e-02  1.892e-02  -4.447 8.99e-06 ***
## Residence_typeUrban 6.466e-04  7.016e-03   0.092 0.926575
## avg_glucose_level 2.898e-04  8.024e-05   3.612 0.000308 ***
## bmi           -5.155e-04  5.197e-04  -0.992 0.321284
## smoking_statusnever smoked -7.336e-03  1.040e-02  -0.705 0.480548
## smoking_statussmokes -5.882e-03  1.240e-02  -0.474 0.635307
## smoking_statusUnknown -1.745e-03  1.163e-02  -0.150 0.880725
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2095 on 3560 degrees of freedom
## Multiple R-squared:  0.09588,    Adjusted R-squared:  0.09181
## F-statistic: 23.59 on 16 and 3560 DF,  p-value: < 2.2e-16
```

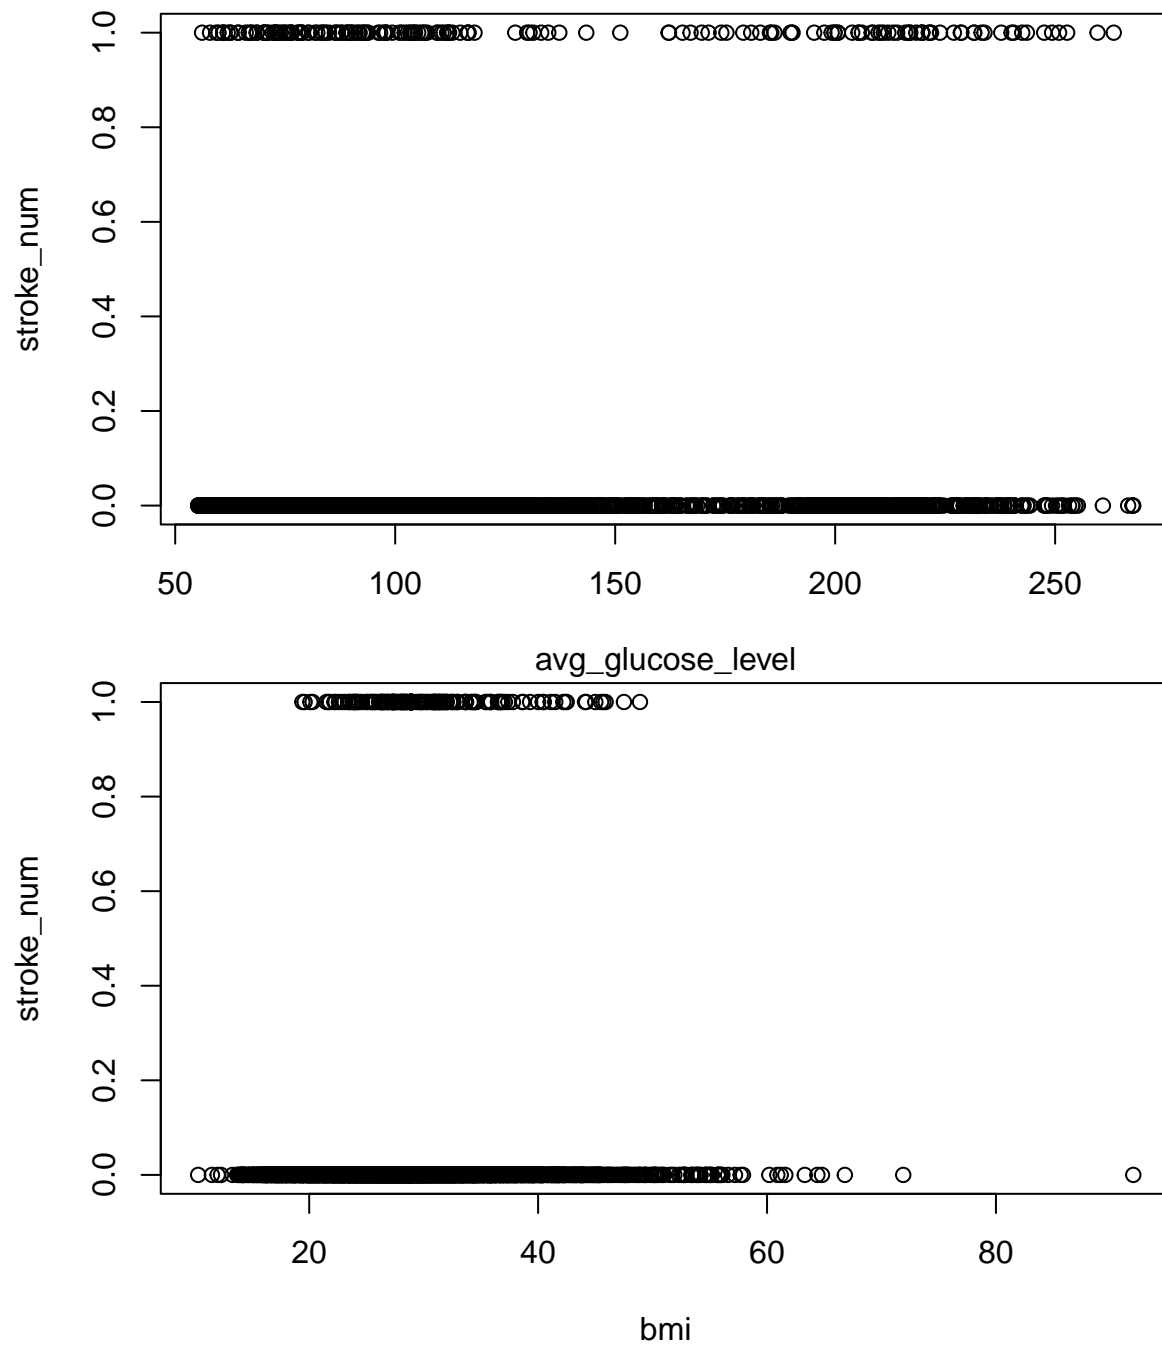
```
plot(stroke_num ~ gender + age + hypertension + heart_disease + ever_married + work_type + Residence_type,
```












```
abline(lm.stroke)
```

```
## Warning in abline(lm.stroke): only using the first two of 17 regression
## coefficients
```



```
# fitting models for simple regression model
```

```
lm.stroke.smoking <- lm(stroke_num ~ smoking_status , data = df_train_linear)
summary(lm.stroke.smoking)
```

```
##
```

```
## Call:
```

```
## lm(formula = stroke_num ~ smoking_status, data = df_train_linear)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -0.08548 -0.05244 -0.04989 -0.03145  0.96855
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.085484   0.008801   9.713 < 2e-16 ***
## smoking_statusnever smoked -0.035597   0.010666  -3.338 0.000854 ***
## smoking_statussmokes    -0.033043   0.012818  -2.578 0.009981 **
## smoking_statusUnknown   -0.054032   0.011040  -4.894 1.03e-06 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```


```
## Residual standard error: 0.2191 on 3573 degrees of freedom
```

```
## Multiple R-squared:  0.006675, Adjusted R-squared:  0.005841
```

```
## F-statistic: 8.004 on 3 and 3573 DF, p-value: 2.579e-05
```

```
plot(stroke_num ~ smoking_status, data = df_train_linear)
abline(lm.stroke.smoking)
```


```
## Warning in abline(lm.stroke.smoking): only using the first two of 4 regression
## coefficients
```



```
# fitting models for simple regression model
lm.stroke.glucose <- lm(stroke_num ~ avg_glucose_level , data = df_train_linear)
summary(lm.stroke.smoking)
```

```
##
## Call:
## lm(formula = stroke_num ~ smoking_status, data = df_train_linear)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.08548 -0.05244 -0.04989 -0.03145  0.96855
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.085484   0.008801    9.713 < 2e-16 ***
## smoking_statusnever smoked -0.035597   0.010666   -3.338 0.000854 ***
## smoking_statussmokes      -0.033043   0.012818   -2.578 0.009981 **
## smoking_statusUnknown     -0.054032   0.011040   -4.894 1.03e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2191 on 3573 degrees of freedom
## Multiple R-squared:  0.006675, Adjusted R-squared:  0.005841
## F-statistic: 8.004 on 3 and 3573 DF, p-value: 2.579e-05

plot(stroke_num ~ avg_glucose_level, data = df_train_linear)
abline(lm.stroke.glucose)
```





```
# fitting models for simple regression model
pred_var <-

lm.stroke.glucose <- lm(stroke_num ~ age + gender , data = df_train_linear)
summary(lm.stroke.smoking)


##
## Call:
## lm(formula = stroke_num ~ smoking_status, data = df_train_linear)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.08548 -0.05244 -0.04989 -0.03145  0.96855
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.085484   0.008801   9.713 < 2e-16 ***
## smoking_statusnever smoked -0.035597   0.010666  -3.338 0.000854 ***
## smoking_statussmokes    -0.033043   0.012818  -2.578 0.009981 **
## smoking_statusUnknown    -0.054032   0.011040  -4.894 1.03e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2191 on 3573 degrees of freedom
## Multiple R-squared:  0.006675, Adjusted R-squared:  0.005841
## F-statistic: 8.004 on 3 and 3573 DF, p-value: 2.579e-05

plot(stroke_num ~ age + gender, data = df_train_linear)
```



```
abline(lm.stroke.glucose)
```


```
## Warning in abline(lm.stroke.glucose): only using the first two of 4 regression
## coefficients
```



```
# fitting models for simple regression model
lm.stroke.bmi <- lm(bmi ~ avg_glucose_level, data = df_train_linear)
summary(lm.stroke.bmi)
```

```
##
## Call:
## lm(formula = bmi ~ avg_glucose_level, data = df_train_linear)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.116  -4.919  -0.754   3.807  64.349
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    26.109983    0.320458  81.477  <2e-16 ***
## avg_glucose_level 0.027086    0.002752   9.843  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.554 on 3575 degrees of freedom
## Multiple R-squared:  0.02639,    Adjusted R-squared:  0.02612
## F-statistic: 96.89 on 1 and 3575 DF,  p-value: < 2.2e-16
plot(bmi ~ avg_glucose_level, data = df_train_linear)
abline(lm.stroke.bmi)
```



5 Generalised Linear Model with family set to Poisson


```
# test
glm.stroke.poisson <- glm(stroke_num ~ age + heart_disease + hypertension + avg_glucose_level,
family = "poisson",
data = df_train_linear)
summary(glm.stroke.poisson)

##
## Call:
## glm(formula = stroke_num ~ age + heart_disease + hypertension +
##      avg_glucose_level, family = "poisson", data = df_train_linear)
##
```

```
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.0558  -0.3142  -0.1712  -0.0831   3.4861
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -7.384780    0.413043 -17.879 < 2e-16 ***
## age           0.066499    0.005875  11.319 < 2e-16 ***
## heart_disease  0.356522    0.186276   1.914 0.05563 .
## hypertension  0.296479    0.168541   1.759 0.07856 .
## avg_glucose_level 0.003448    0.001222   2.823 0.00476 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 1084.09  on 3576  degrees of freedom
## Residual deviance:  795.51  on 3572  degrees of freedom
## AIC: 1169.5
##
## Number of Fisher Scoring iterations: 7

#plot(stroke_num ~ age + heart_disease + hypertension + avg_glucose_level, data = df_train_linear)
#abline(glm.stroke_poisson)
ggplot(data = df_train_linear, aes(x = avg_glucose_level, y = stroke_num)) +
  geom_jitter(width = 0, height = 0.05) +
  geom_smooth(method = "glm", method.args = list(family = "poisson"))

## `geom_smooth()` using formula 'y ~ x'
```



```
# Evaluating model fit poisson
```

```
# fitted(glm.stroke.poisson)
```

```
fitted.glm.stroke.poisson <- ifelse(fitted(glm.stroke.poisson) < 0.08, yes = 0, no = 1)
```

```
head(fitted.glm.stroke.poisson)
```

```
## 1 2 3 4 5 6
```

```
## 1 1 0 1 1 1
```

```
obs.fitted.comp.poisson <- data.frame(obs = df_train_linear$stroke_num, fitted = fitted.glm.stroke.poisson)
```

```
table(obs = obs.fitted.comp.poisson$obs, fit = obs.fitted.comp.poisson$fitted)
```

```
##      fit
```

```
## obs    0    1
```

```
##    0 2809  586
```

```
##    1   51  131
```

```
table(obs = obs.fitted.comp.poisson$obs, fit = obs.fitted.comp.poisson$fitted) %>%  
  prop.table() %>%  
  round(digits = 2)
```

```
##      fit
```

```
## obs    0    1
```

```
##    0 0.79 0.16
```

```
##    1 0.01 0.04
```

6 Generalised Linear Model with family set to Binomial

Since we are essentially dealing with a classification issue, using logistic regression in the form of a GLM with family set to “binomial” is the best method to apply out of all the models introduced so far. For this reason, we shall go into more detail here.


```
# Include all variables to start variable selection
# Plot all of them for visual analysis
glm.stroke.binomial <- glm(stroke ~ .,
family = "binomial",
data = df_train_linear)


## Warning: glm.fit: algorithm did not converge

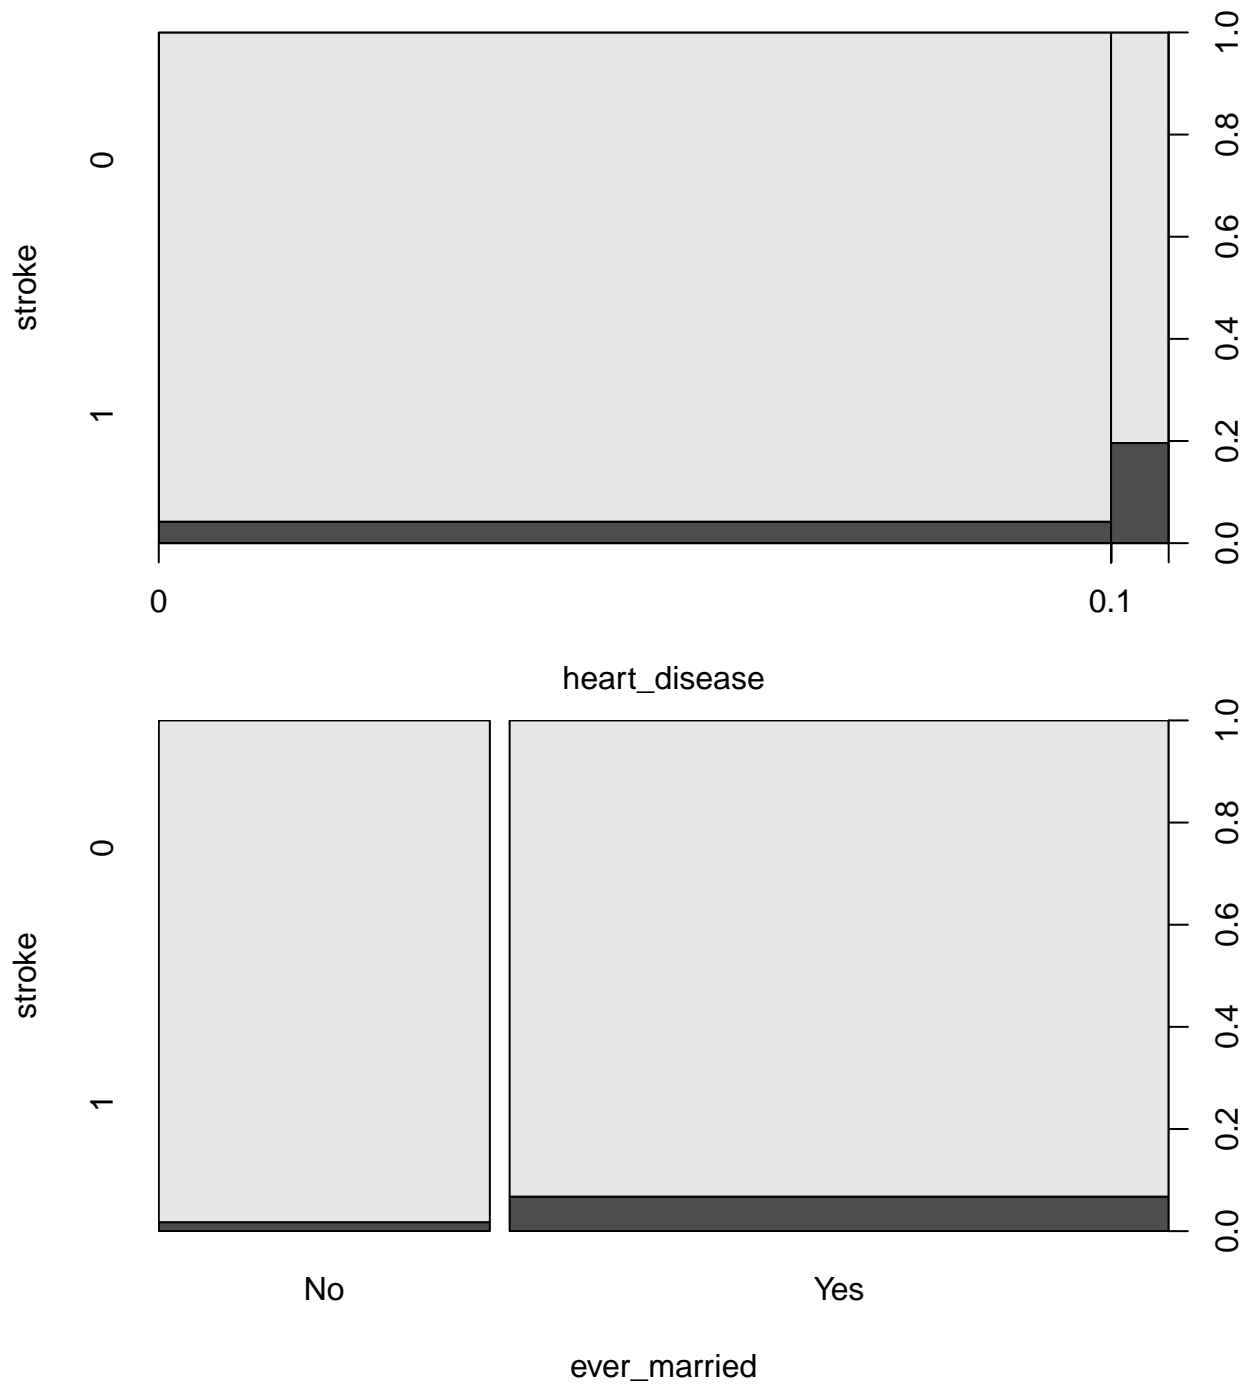
summary(glm.stroke.binomial)


##
## Call:
## glm(formula = stroke ~ ., family = "binomial", data = df_train_linear)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.409e-06 -2.409e-06 -2.409e-06 -2.409e-06  2.409e-06
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -2.657e+01  3.576e+04  -0.001    0.999
## id              3.753e-17  2.789e-01   0.000    1.000
## genderMale     -1.296e-12  1.228e+04   0.000    1.000
## genderOther    -1.930e-12  3.569e+05   0.000    1.000
## age            -5.937e-13  4.548e+02   0.000    1.000
## hypertension   -1.935e-11  2.060e+04   0.000    1.000
## heart_disease   2.369e-10  2.725e+04   0.000    1.000
## ever_marriedYes  2.975e-13  1.780e+04   0.000    1.000
## work_typeGovt_job  1.343e-11  3.160e+04   0.000    1.000
## work_typeNever_worked  4.315e-12  9.134e+04   0.000    1.000
## work_typePrivate  1.297e-11  2.654e+04   0.000    1.000
## work_typeSelf-employed  1.753e-12  3.226e+04   0.000    1.000
## Residence_typeUrban  7.404e-12  1.193e+04   0.000    1.000
## avg_glucose_level  1.803e-13  1.367e+02   0.000    1.000
## bmi             9.729e-14  8.838e+02   0.000    1.000
## smoking_statusnever smoked  4.215e-12  1.768e+04   0.000    1.000
## smoking_statussmokes  1.480e-11  2.109e+04   0.000    1.000
## smoking_statusUnknown  6.158e-12  1.978e+04   0.000    1.000
## stroke_num       5.313e+01  2.850e+04   0.002    0.999
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1.4387e+03  on 3576  degrees of freedom
## Residual deviance: 2.0752e-08  on 3558  degrees of freedom
## AIC: 38
##
## Number of Fisher Scoring iterations: 25


plot(stroke ~ ., data = df_train_linear)
```















```
abline(glm.stroke.binomial)
```


```
## Warning in abline(glm.stroke.binomial): only using the first two of 19
## regression coefficients
```

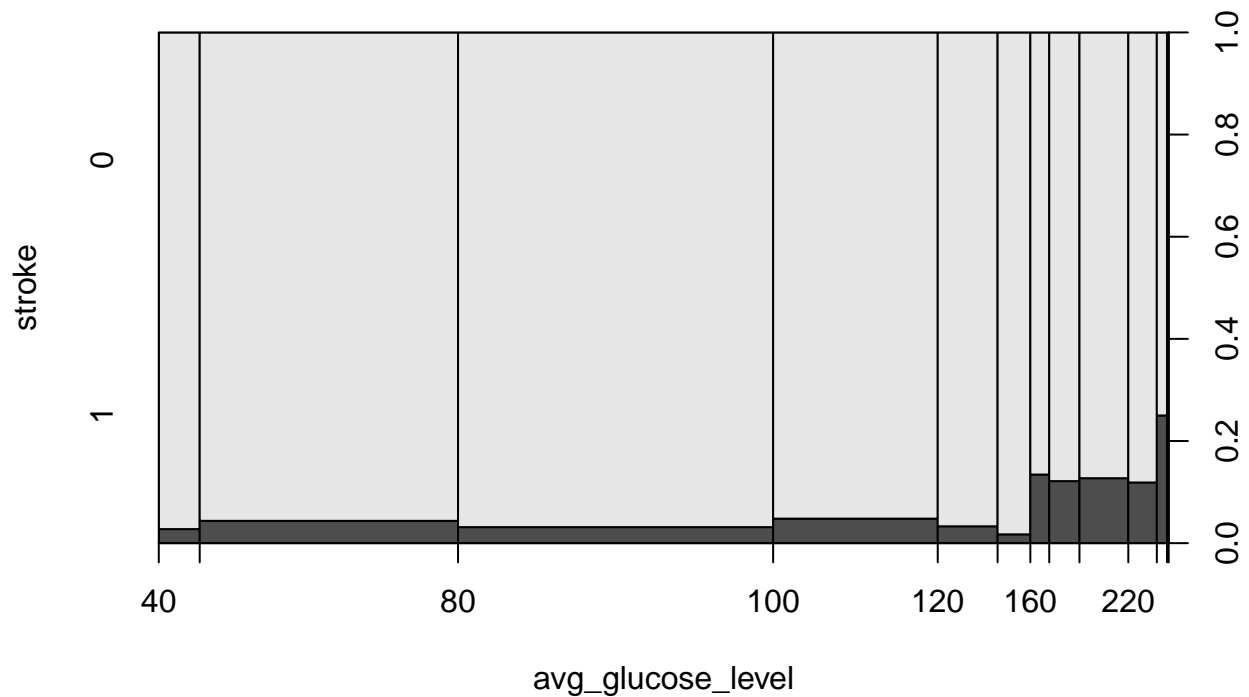


```
# First iteration with parameters chosen from intuitive domain knowledge and exploratory analysis of data
glm.stroke.binomial <- glm(stroke ~ age + heart_disease + hypertension + work_type + avg_glucose_level +
family = "binomial",
data = df_train_linear)
summary(glm.stroke.binomial)
```

```
##
## Call:
## glm(formula = stroke ~ age + heart_disease + hypertension + work_type +
##      avg_glucose_level + smoking_status, family = "binomial",
##      data = df_train_linear)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.1696  -0.3241  -0.1595  -0.0771   3.6785
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -7.037953   1.036288  -6.792 1.11e-11 ***
## age              0.078346   0.006958  11.260 < 2e-16 ***
## heart_disease    0.431469   0.213875   2.017  0.0437 *
## hypertension     0.378996   0.189527   2.000  0.0455 *
## work_typeGovt_job -1.003706   1.105068  -0.908  0.3637
## work_typeNever_worked -10.077332 362.934034  -0.028  0.9778
## work_typePrivate  -0.818567   1.088363  -0.752  0.4520
## work_typeSelf-employed -1.382649   1.112789  -1.243  0.2140
## avg_glucose_level  0.003936   0.001360   2.894  0.0038 **
## smoking_statusnever smoked -0.186796   0.203452  -0.918  0.3586
## smoking_statussmokes  0.032194   0.256225   0.126  0.9000
## smoking_statusUnknown -0.081986   0.243356  -0.337  0.7362
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1438.7  on 3576  degrees of freedom
## Residual deviance: 1117.5  on 3565  degrees of freedom
## AIC: 1141.5
##
## Number of Fisher Scoring iterations: 14
plot(stroke ~ age + heart_disease + hypertension + work_type + avg_glucose_level + smoking_status, data = df_train_linear)
```








```
abline(glm.stroke.binomial)
```

```
## Warning in abline(glm.stroke.binomial): only using the first two of 12
## regression coefficients
```



```
# second iteration only keeping statistically relevant parameters from previous model
glm.stroke.binomial.2 <- glm(stroke_num ~ age + heart_disease + hypertension + avg_glucose_level,
family = "binomial",
data = df_train_linear)
summary(glm.stroke.binomial.2)
```




```
##
## Call:
## glm(formula = stroke_num ~ age + heart_disease + hypertension +
##      avg_glucose_level, family = "binomial", data = df_train_linear)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.1183  -0.3237  -0.1671  -0.0764   3.8413
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -7.739236   0.443953 -17.433  <2e-16 ***
## age           0.072374   0.006312  11.467  <2e-16 ***
## heart_disease  0.465121   0.211312   2.201   0.0277 *
## hypertension  0.362843   0.187103   1.939   0.0525 .
## avg_glucose_level 0.004117   0.001351   3.048   0.0023 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1438.7  on 3576  degrees of freedom
## Residual deviance: 1127.6  on 3572  degrees of freedom
## AIC: 1137.6
##
## Number of Fisher Scoring iterations: 7

# plot(stroke ~ age + heart_disease + hypertension + avg_glucose_level, data = df_train_linear)
# abline(glm.stroke.binomial.2)

pred <- predict(glm.stroke.binomial.2)
# pred

ggplot(data = df_train_linear, aes(x = age, y = stroke_num)) +
  geom_jitter(width = 0, height = 0.05) +
  geom_smooth(method = "glm", method.args = list(family = "binomial"))

## `geom_smooth()` using formula 'y ~ x'
```



```
# Evaluating model fit using fitted()
```

```
# fitted(glm.stroke.binomial.2)
```

```
fitted.glm.stroke.binomial.2 <- ifelse(fitted(glm.stroke.binomial.2) < 0.2, yes = 0, no = 1)
head(fitted.glm.stroke.binomial.2)
```

```
## 1 2 3 4 5 6
## 0 1 0 1 1 1
```

```
obs.fitted.comp <- data.frame(obs = df_train_linear$stroke_num, fitted = fitted.glm.stroke.binomial.2)
```

```
table(obs = obs.fitted.comp$obs, fit = obs.fitted.comp$fitted)
```

```
##      fit
## obs    0    1
##    0 3235 160
##    1  130  52
```

```
table(obs = obs.fitted.comp$obs, fit = obs.fitted.comp$fitted) %>%
  prop.table() %>%
  round(digits = 2)
```

```
##      fit
## obs    0    1
##    0 0.90 0.04
##    1 0.04 0.01
```

```
# Evaluating model fit using predict
```

```
# predict(glm.stroke.binomial.2, df_test_linear, type = "response")

# predicted(glm.stroke.binomial.2)
predicted.glm.stroke.binomial.2 <- ifelse(predict(glm.stroke.binomial.2, df_test_linear, type = "response"), 1, 0)
head(predicted.glm.stroke.binomial.2)

## 1 2 3 4 5 6
## 0 0 0 0 0 0

obs.predicted.comp <- data.frame(obs = df_test_linear$stroke_num, predicted = predicted.glm.stroke.binomial.2)

table(obs = obs.predicted.comp$obs, fit = obs.predicted.comp$predicted)

##      fit
## obs    0    1
##   0 1412   54
##   1   54   13

table(obs = obs.predicted.comp$obs, fit = obs.predicted.comp$predicted) %>%
  prop.table() %>%
  round(digits = 2)

##      fit
## obs    0    1
##   0 0.92 0.04
##   1 0.04 0.01
```

7 Generalised Additive Model


```
library(mgcv)

gam.stroke <- gam(stroke ~ smoking_status + s(bmi),
  family = "binomial",
  data = df_train_linear)
summary(gam.stroke)

##
## Family: binomial
## Link function: logit
##
## Formula:
## stroke ~ smoking_status + s(bmi)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -2.6671     0.1636 -16.302 < 2e-16 ***
## smoking_statusnever smoked  -0.5438     0.1924  -2.827  0.00470 **
## smoking_statussmokes        -0.5377     0.2398  -2.243  0.02492 *
## smoking_statusUnknown       -0.7270     0.2298  -3.164  0.00155 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq p-value
```

```
## s(bmi) 4.745  5.588  26.22 0.000161 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.0132   Deviance explained =  4.5%
## UBRE = -0.61102   Scale est. = 1          n = 3577
plot(stroke ~ smoking_status, data = df_train_linear)
abline(gam.stroke)

## Warning in abline(gam.stroke): only using the first two of 13 regression
## coefficients
```



8 Neural Network Yves

9 Support Vector Machine (Larissa)

Stroke Data Classification using a Support Vector Machine.

```
ytrain = df_train$stroke
ytest = df_test$stroke


svm_model <- svm(stroke ~ ., data = df_train, type = "C-classification", kernel = "radial", cost = "5")
summary(svm_model)

##
## Call:
## svm(formula = stroke ~ ., data = df_train, type = "C-classification",
##      kernel = "radial", cost = "5")
##
##
## Parameters:
##   SVM-Type:  C-classification
```

```
## SVM-Kernel: radial
##      cost: 5
##
## Number of Support Vectors: 95
##
## ( 31 64 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

```
plot(svm_model, data = df_train, bmi ~ age, slice = list(avg_glucose_level = 3))
```

SVM classification plot



```
#sum_training_prediction <- predict(svm_model, newdata = df_train)
#sum_training_error <- mean(sum_training_prediction != ytrain)
#draw_confusion_matrix(confusionMatrix(sum_training_prediction, df_train$stroke), "Stroke", "No Stroke")
#sum_training_error
#confusionMatrix(sum_training_prediction, df_train$stroke)

#sum_prediction <- predict(svm_model, newdata = df_test)
#sum_test_error <- mean(sum_prediction != ytest)
#draw_confusion_matrix(confusionMatrix(sum_prediction, df_test$stroke), "Stroke", "No Stroke")
#sum_test_error
```

10 OPTIONAL solve an optimisation problem