# Applied Machine Learning and Predictive Modelling I: Modelling Stroke Data

Authors: Larissa Eisele, Fabian Lüthard, Yves Maillard

Module: Applied Machine Learning and Predictive Modelling I

Submitted on 10th of June, 2022

SUPERVISOR: MATTEO TANADINI AND DANIEL MEISTER

Lucerne University of Applied Sciences and Arts

# Table of Contents

1	Introduction	1
<b>2</b>	Importing Data	1
	2.1 Summaries	2
	2.2 Scatterplots	3
	2.3 Boxplots	7
3	Methodology	13
4	Linear Model	<b>1</b> 4
5	Generalised Linear Model with family set to Binomial Fabian	16
6	Generalised Additive Model Fabian	16
7	Neural Network Yves	16
8	Support Vector Machine (Larissa)	16
9	OPTIONAL solve an optimisation problem	18

#### 1 Introduction

Use case: We are a smart watch manufacturer working on a new feature for stroke prevention. We are going to analyze survey data that we plan to ask our users, complementing it with HR (Heart Rate) and CGM (Continuous Glucose Monitoring) data that our product already measures. We hope that our feature can prevent serious health issues and motivate our users to adopt healthier lifestyles.

### 2 Importing Data

```
stroke_data <- read_csv('./data/healthcare-dataset-stroke-data.csv')</pre>
## Rows: 5110 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (6): gender, ever_married, work_type, Residence_type, bmi, smoking_status
## dbl (6): id, age, hypertension, heart_disease, avg_glucose_level, stroke
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
s <- spec(stroke_data)</pre>
## cols(
     id = col_double(),
     gender = col_character(),
##
##
     age = col_double(),
##
    hypertension = col_double(),
##
    heart_disease = col_double(),
##
     ever_married = col_character(),
##
     work_type = col_character(),
##
     Residence_type = col_character(),
##
     avg_glucose_level = col_double(),
##
     bmi = col_character(),
     smoking_status = col_character(),
##
##
     stroke = col_double()
## )
cols_condense(s)
## cols(
##
     .default = col_character(),
     id = col_double(),
##
     age = col_double(),
    hypertension = col_double(),
##
##
    heart_disease = col_double(),
     avg_glucose_level = col_double(),
     stroke = col_double()
##
## )
```

```
stroke_data$bmi <- as.numeric(stroke_data$bmi)</pre>
## Warning: NAs introduced by coercion
stroke_data$bmi[is.na(stroke_data$bmi)] <- mean(stroke_data$bmi, na.rm = TRUE)</pre>
stroke_data$stroke <- as.factor(stroke_data$stroke)</pre>
stroke_data$age <- as.integer(stroke_data$age)</pre>
stroke_data$smoking_status <- as.factor(stroke_data$smoking_status)</pre>
stroke_data$work_type <- as.factor(stroke_data$work_type)</pre>
stroke_data$gender <- as.factor(stroke_data$gender)</pre>
stroke_data$ever_married <- as.factor(stroke_data$ever_married)</pre>
stroke_data$Residence_type <- as.factor(stroke_data$Residence_type)</pre>
stroke_data
## # A tibble: 5,110 x 12
##
        id gender age hypertension heart_disease ever_married work_type
                                          <dbl> <fct>
##
     <dbl> <fct> <int> <dbl>
                                                              <fct>
## 1 9046 Male
                    67
                                 0
                                               1 Yes
                                                              Private
## 2 51676 Female 61
                                  0
                                                0 Yes
                                                              Self-employed
## 3 31112 Male 80
                                  0
                                                1 Yes
                                                              Private
## 4 60182 Female 49
                                 0
                                                0 Yes
                                                             Private
## 5 1665 Female 79
                                 1
                                                             Self-employed
                                                0 Yes
## 6 56669 Male 81
                                 0
                                                0 Yes
                                                              Private
                  74
## 7 53882 Male
                                 1
                                               1 Yes
                                                             Private
## 8 10434 Female 69
                                 0
                                                O No
                                                             Private
## 9 27419 Female 59
                                  0
                                                0 Yes
                                                             Private
## 10 60491 Female 78
                                  0
                                                0 Yes
                                                              Private
## # ... with 5,100 more rows, and 5 more variables: Residence_type <fct>,
## # avg_glucose_level <dbl>, bmi <dbl>, smoking_status <fct>, stroke <fct>
2.1
     Summaries
#testing the effect of non smokers and smokers
count_by_smoke_status <- stroke_data %>%
 select(smoking_status, stroke) %>%
 group_by(smoking_status, stroke) %>%
 summarise(N = n())
```

```
## 'summarise()' has grouped output by 'smoking_status'. You can override using the
## '.groups' argument.

#testing the effect of work type
count_by_work_type <- stroke_data %>%
    select(work_type, stroke) %>%
    group_by(work_type, stroke) %>%
    summarise( N = n())

## 'summarise()' has grouped output by 'work type'. You can override using the
```

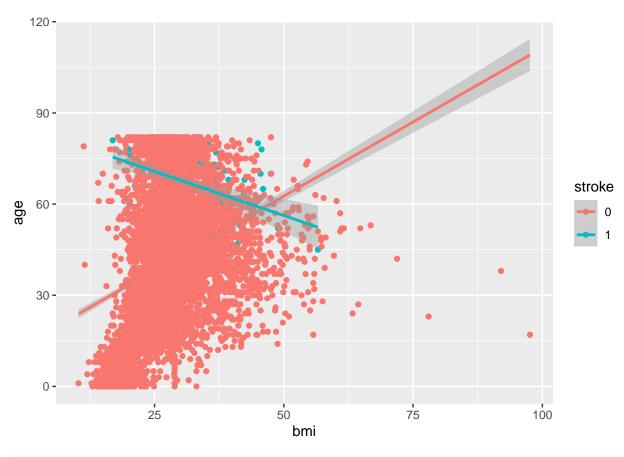
## '.groups' argument.

```
# testing the effects of residence type
count_by_Residence_type <- stroke_data %>%
   select(Residence_type, stroke) %>%
   group_by(Residence_type, stroke) %>%
  summarise(N = n())
## 'summarise()' has grouped output by 'Residence_type'. You can override using the
## '.groups' argument.
# testing the effects of gender
count_by_gender <- stroke_data %>%
   select(gender, stroke) %>%
   group_by(gender, stroke) %>%
  summarise(N = n())
## 'summarise()' has grouped output by 'gender'. You can override using the
## '.groups' argument.
# testing the effects of hypertension
count_by_hypertension <- stroke_data %>%
   select(hypertension, stroke) %>%
   group_by(hypertension, stroke) %>%
  summarise(N = n())
## 'summarise()' has grouped output by 'hypertension'. You can override using the
## '.groups' argument.
# testing the effects of heart disease
count_by_heart_disease <- stroke_data %>%
   select(heart_disease, stroke) %>%
  group_by(heart_disease, stroke) %>%
  summarise(N = n())
## 'summarise()' has grouped output by 'heart_disease'. You can override using the
## '.groups' argument.
# testing the effects of marriage status
count_by_marriage <- stroke_data %>%
   select(ever_married, stroke) %>%
  group_by(ever_married, stroke) %>%
  summarise(N = n())
## 'summarise()' has grouped output by 'ever_married'. You can override using the
## '.groups' argument.
```

#### 2.2 Scatterplots

```
stroke_data %>%
  ggplot(mapping = aes(x = bmi, y = age, color = stroke)) +
  geom_point() +
  geom_smooth(method = 'lm')
```

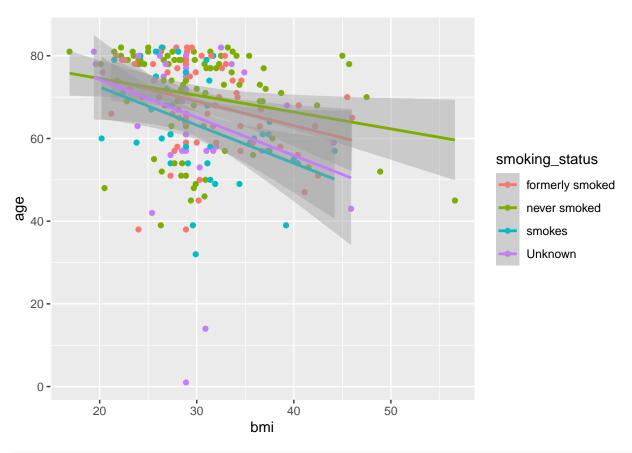
## 'geom\_smooth()' using formula 'y ~ x'



```
stroke_data %>%
  filter(stroke == 1) %>%
  ggplot(mapping = aes(x = bmi, y = age, color = smoking_status)) +
  geom_point(method = 'lm') +
  geom_smooth(method = 'lm')
```

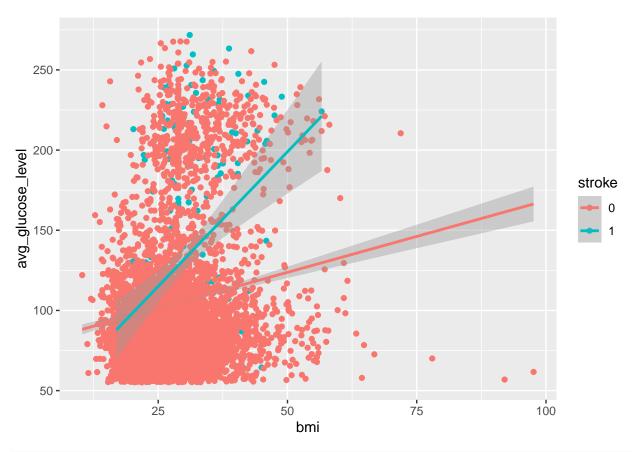
## Warning: Ignoring unknown parameters: method

## 'geom\_smooth()' using formula 'y ~ x'



```
stroke_data %>%
  ggplot(mapping = aes(x = bmi, y = avg_glucose_level, color = stroke)) +
  geom_point() +
  geom_smooth(method = 'lm')
```

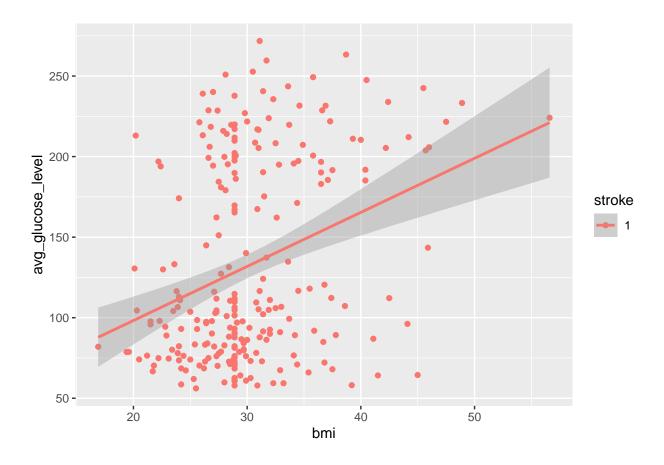
## 'geom\_smooth()' using formula 'y ~ x'



```
stroke_data %>%
  filter(stroke == 1) %>%
  ggplot(mapping = aes(x = bmi, y = avg_glucose_level, color = stroke)) +
  geom_point(method = 'lm') +
  geom_smooth(method = 'lm')
```

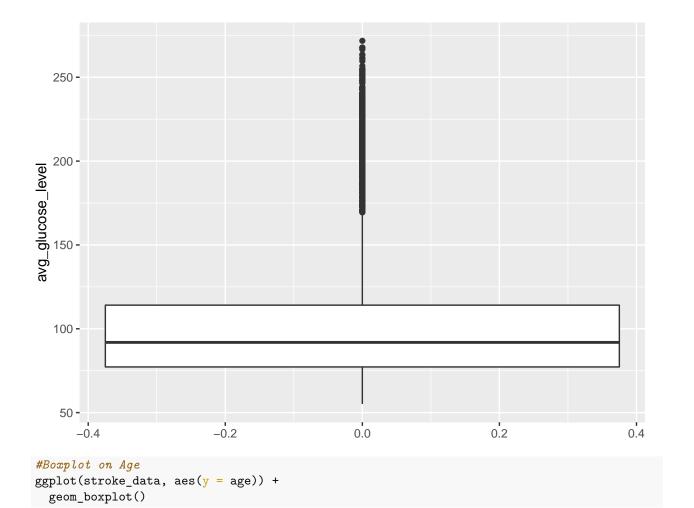
## Warning: Ignoring unknown parameters: method

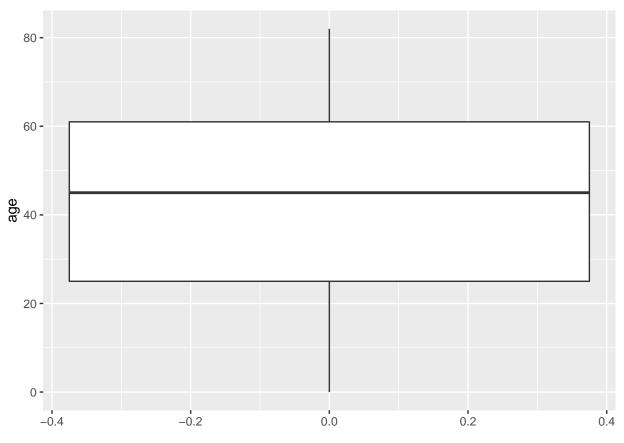
## 'geom\_smooth()' using formula 'y ~ x'

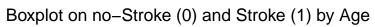


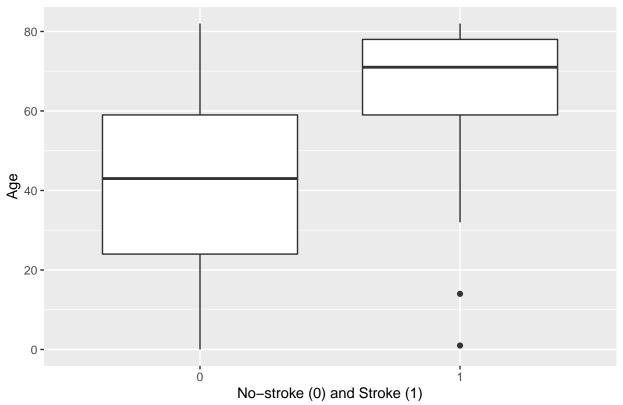
## 2.3 Boxplots

```
#Boxplot on avg. glucose_level
ggplot(stroke_data, aes(y = avg_glucose_level)) +
  geom_boxplot()
```

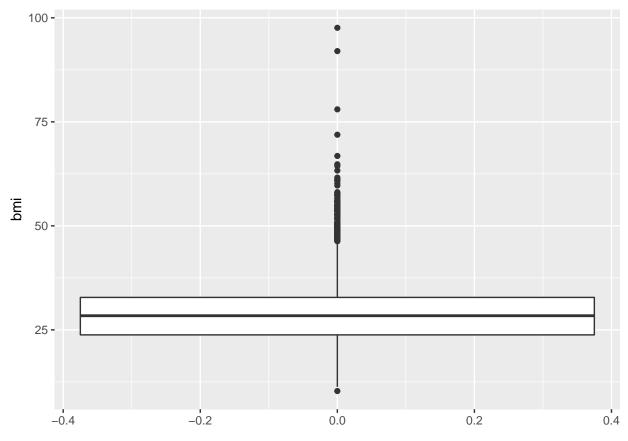


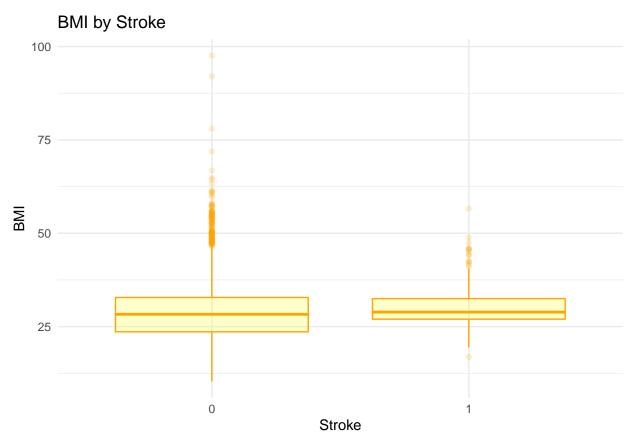


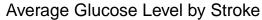


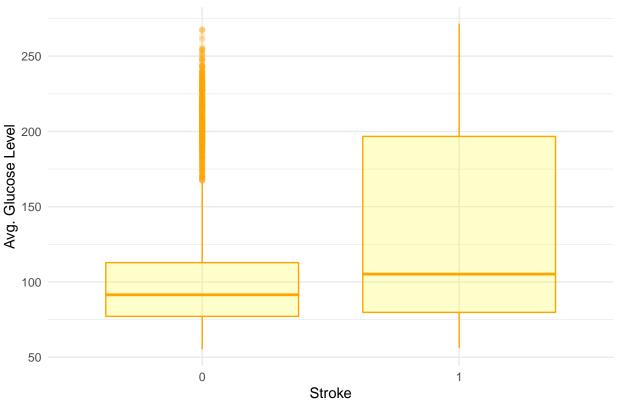


```
#Boxplot on BMI
ggplot(stroke_data, aes(y = bmi)) +
geom_boxplot()
```









## 3 Methodology

## 1 51676 Female

## 2 10434 Female

## 3 60491 Female

61

78

0

0

0

```
set.seed(7406)
n=dim(stroke_data[1]) # number of observations in dataset
n_train=0.70*n # training set is 70%
flag = sort(sample(1:n, size=n_train, replace=FALSE))
## Warning in 1:n: numerical expression has 2 elements: only the first used
\# Use df (all data points without ID column) df_train, and df_test
# Gender, hypertension, heart disease, ever married, work type, residence type, smoking status, and str
# This should allow for the best modeling options possible for our methods.
df_train = stroke_data[flag,]
df_test = stroke_data[-flag,]
head(df_test)
## # A tibble: 6 x 12
                    age hypertension heart_disease ever_married work_type
##
        id gender
     <dbl> <fct> <int>
                               <dbl>
                                             <dbl> <fct>
```

0 Yes

O No

0 Yes

Self-employed

Private

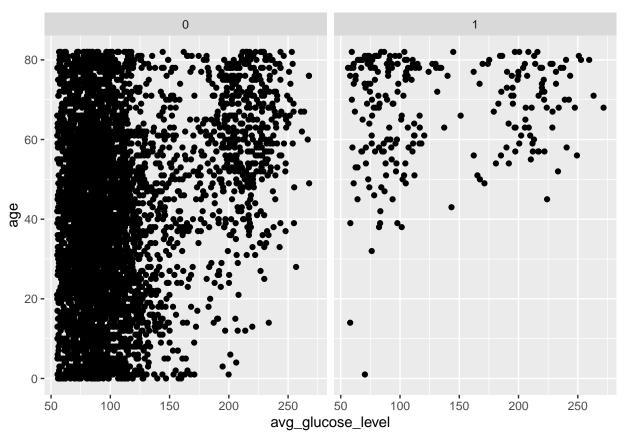
Private

```
## 4 12095 Female
                                                 1 Yes
                                                                Govt_job
                     61
## 5 58202 Female
                     50
                                   1
                                                 0 Yes
                                                                Self-employed
## 6 56112 Male
                     64
                                   0
                                                 1 Yes
                                                                Private
## # ... with 5 more variables: Residence_type <fct>, avg_glucose_level <dbl>,
## # bmi <dbl>, smoking_status <fct>, stroke <fct>
```

head(df\_train)

```
## # A tibble: 6 x 12
                   age hypertension heart_disease ever_married work_type
##
       id gender
    <dbl> <fct> <int>
                              <dbl>
                                            <dbl> <fct>
##
                                                               <fct>
## 1 9046 Male
                                                1 Yes
                    67
                                  0
                                                               Private
## 2 31112 Male
                    80
                                  0
                                                1 Yes
                                                               Private
## 3 60182 Female
                    49
                                  0
                                                0 Yes
                                                               Private
## 4 1665 Female
                    79
                                  1
                                                0 Yes
                                                               Self-employed
## 5 56669 Male
                    81
                                  0
                                                0 Yes
                                                               Private
                    74
## 6 53882 Male
                                  1
                                                1 Yes
                                                               Private
## # ... with 5 more variables: Residence_type <fct>, avg_glucose_level <dbl>,
## # bmi <dbl>, smoking_status <fct>, stroke <fct>
```

#### 4 Linear Model



```
# fitting models for simple regression model
lm.stroke <- lm(age ~ avg_glucose_level, data = stroke_data)
summary(lm.stroke)</pre>
```

```
##
## Call:
## lm(formula = age ~ avg_glucose_level, data = stroke_data)
##
## Residuals:
      Min
               1Q Median
                               ЗQ
##
## -53.362 -16.762
                    1.127 16.638 44.662
##
## Coefficients:
##
                     Estimate Std. Error t value Pr(>|t|)
## (Intercept)
                    30.584911
                                0.783855
                                           39.02
                                                   <2e-16 ***
## avg_glucose_level 0.118988
                                0.006792
                                           17.52
                                                   <2e-16 ***
## ---
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' 1
##
## Residual standard error: 21.99 on 5108 degrees of freedom
## Multiple R-squared: 0.05667, Adjusted R-squared: 0.05649
## F-statistic: 306.9 on 1 and 5108 DF, p-value: < 2.2e-16
```

# Generalised Linear Model with family set to Poisson

""r

```
glm.stroke <- glm(age ~ smoking_status,</pre>
family = "poisson",
data = stroke_data)
summary(glm.stroke)
##
## Call:
## glm(formula = age ~ smoking_status, family = "poisson", data = stroke_data)
## Deviance Residuals:
      Min
##
                     Median
                                   3Q
                 1Q
                                           Max
## -7.7708 -3.0398 -0.1092
                               2.1971
                                        7.7616
##
## Coefficients:
##
                               Estimate Std. Error z value Pr(>|z|)
                              4.006059
                                         0.004535 883.27
                                                             <2e-16 ***
## (Intercept)
## smoking_statusnever smoked -0.161357
                                          0.005646 -28.58
                                                             <2e-16 ***
## smoking_statussmokes
                                          0.006891 -22.33
                                                             <2e-16 ***
                             -0.153864
                                         0.006482 -92.32
## smoking_statusUnknown
                              -0.598470
                                                             <2e-16 ***
## ---
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
      Null deviance: 72137 on 5109 degrees of freedom
## Residual deviance: 61972 on 5106 degrees of freedom
## AIC: 89244
##
## Number of Fisher Scoring iterations: 5
#glm.stroke <- glm(bmi ~ stroke,
#family = "poisson",
\#data = stroke_data
#summary(qlm.stroke)
```

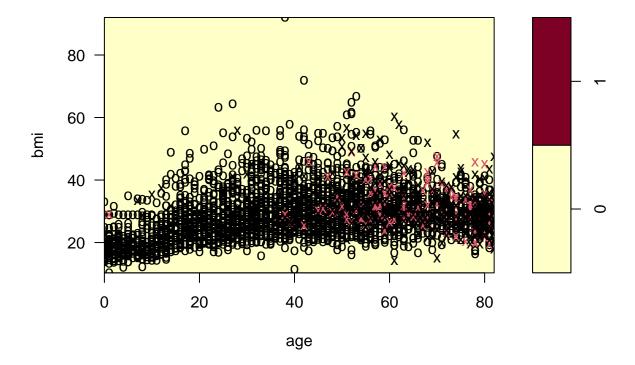
- 5 Generalised Linear Model with family set to Binomial Fabian
- 6 Generalised Additive Model Fabian
- 7 Neural Network Yves
- 8 Support Vector Machine (Larissa)

Stroke Data Classification using a Support Vector Machine.

```
ytrain = df_train$stroke
ytest = df_test$stroke
```

```
svm_model <- svm(stroke ~. , data = df_train, type = "C-classification", kernel = "radial", cost = "5")</pre>
summary(svm_model)
##
## Call:
## svm(formula = stroke ~ ., data = df_train, type = "C-classification",
##
       kernel = "radial", cost = "5")
##
##
## Parameters:
##
      SVM-Type: C-classification
    SVM-Kernel: radial
##
##
          cost: 5
##
## Number of Support Vectors: 600
##
    ( 182 418 )
##
##
##
## Number of Classes: 2
##
## Levels:
   0 1
##
plot(svm_model, data = df_train, bmi ~ age, slice = list(avg_glucose_level = 3))
```

## **SVM** classification plot



```
#sum_training_prediction <- predict(sum_model, newdata = df_train)
#sum_training_error <- mean(sum_training_prediction != ytrain)
#draw_confusion_matrix(confusionMatrix(sum_training_prediction,df_train$stroke), "Stroke", "No Stroke")
#sum_training_error
#confusionMatrix(sum_training_prediction,df_train$stroke)

#sum_prediction <- predict(sum_model, newdata = df_test)
#sum_test_error <- mean(sum_prediction != ytest)
#draw_confusion_matrix(confusionMatrix(sum_prediction,df_test$stroke), "Stroke", "No Stroke")
#sum_test_error</pre>
```

## 9 OPTIONAL solve an optimisation problem